

Healthcare Claims Database - Installation Guide for DBAs

Prerequisites

Required Privileges

- ACCOUNTADMIN or SECURITYADMIN role for security setup
- SYSADMIN role for database and schema creation
- CREATE DATABASE privilege
- CREATE WAREHOUSE privilege

System Requirements

- Snowflake Enterprise Edition (recommended for resource monitors)
- Minimum 100 credit allocation for demo environment
- Network access for demo users

Installation Steps

Step 1: Security and Access Setup

Script: `snowflake_security_setup.sql`

Run as: SECURITYADMIN or ACCOUNTADMIN

Duration: ~2 minutes

```
sql

-- Execute the entire security setup script
-- This creates roles, users, warehouse, and resource monitors
```

Validates:

- Warehouse `healthcare_demo_wh` created
- 5 roles created (analyst, power_analyst, report_writer, data_steward, admin)
- 6 demo users created with passwords
- Resource monitor configured (100 credits/month)

Verification:

```
sql
```

```
SHOW WAREHOUSES LIKE 'healthcare_demo_wh';  
SHOW ROLES LIKE 'healthcare%';  
SHOW USERS LIKE 'demo%';
```

Step 2: Database and Schema Creation

Script: `healthcare_claims_ddl.sql`

Run as: SYSADMIN or role with CREATE DATABASE privilege

Duration: ~1 minute

```
sql  
  
-- Create database if not exists  
CREATE DATABASE IF NOT EXISTS healthcare_db;  
USE DATABASE healthcare_db;  
  
-- Execute DDL script to create all tables  
-- Creates 17 tables with foreign keys and indexes
```

Validates:

- Database `healthcare_db` exists
- Schema `healthcare` created
- 17 tables created with proper relationships
- Indexes created for performance

Verification:

```
sql  
  
USE DATABASE healthcare_db;  
USE SCHEMA healthcare;  
SHOW TABLES;  
SELECT COUNT(*) FROM INFORMATION_SCHEMA.TABLES  
WHERE TABLE_SCHEMA = 'HEALTHCARE';  
-- Expected: 17 tables
```

Step 2a: RAW Schema Creation (For ETL/dbt Pipeline)

Script: `raw_schema_ddl.sql`

Run as: SYSADMIN

Duration: ~1 minute

Purpose: Creates landing zone for CSV data before transformation

```
sql
```

```
USE DATABASE healthcare_db;
```

```
-- Execute raw_schema_ddl.sql to create RAW schema with 17 staging tables
```

Key Differences from Final Schema:

- All columns are VARCHAR (no type constraints)
- No foreign keys (allows any order loading)
- Includes metadata columns: `_loaded_at`, `_file_name`
- CSV stage created with appropriate formatting

Loading Process:

1. Upload CSVs to stage: `PUT file://*.csv @raw.csv_stage`
2. COPY INTO raw tables (no FK constraints)
3. Run dbt to transform RAW → HEALTHCARE schema

Verification:

```
sql
```

```
SHOW TABLES IN SCHEMA raw;
```

```
-- Expected: 17 tables
```

```
SHOW STAGES IN SCHEMA raw;
```

```
-- Expected: csv_stage
```

Step 3: Load Sample Data

Script: `healthcare_sample_data.sql`

Run as: Role with INSERT privilege (healthcare_demo_admin_role)

Duration: ~5 minutes

```
sql
```

```
USE DATABASE healthcare_db;
USE SCHEMA healthcare;

-- Execute sample data script
-- Loads 100 claims with all related records
```

Validates:

- 20 patients loaded
- 100 claims with line items
- Payments, denials, appeals properly linked
- All foreign key relationships satisfied

Verification:

```
sql

-- Run the verification query at end of sample data script
-- Should show counts for all 17 tables
SELECT 'Claims' as entity, COUNT(*) FROM claim
UNION ALL
SELECT 'Claim Line Items', COUNT(*) FROM claim_line_item;
-- Expected: 100 claims, ~400 line items
```

Step 4: Create Reporting Views

Script: healthcare_reporting_views.sql

Run as: Role with CREATE VIEW privilege

Duration: ~1 minute

```
sql

USE DATABASE healthcare_db;
USE SCHEMA healthcare;

-- Execute views script
-- Creates 15 analytical views
```

Validates:

- 15 views created successfully

- Views accessible to analyst roles
- No errors in view definitions

Verification:

```
sql

SHOW VIEWS IN SCHEMA healthcare;
SELECT COUNT(*) FROM INFORMATION_SCHEMA.VIEWS
WHERE TABLE_SCHEMA = 'HEALTHCARE';
-- Expected: 15 views

-- Test a view
SELECT * FROM v_claims_summary_dashboard LIMIT 10;
```

Step 5: Grant Permissions

Run as: SECURITYADMIN

```
sql

-- Verify grants are in place
SHOW GRANTS TO ROLE healthcare_analyst_role;
SHOW GRANTS TO ROLE healthcare_power_analyst_role;

-- Grant future object privileges if needed
GRANT SELECT ON FUTURE TABLES IN SCHEMA healthcare_db.healthcare
TO ROLE healthcare_analyst_role;
GRANT SELECT ON FUTURE VIEWS IN SCHEMA healthcare_db.healthcare
TO ROLE healthcare_analyst_role;
```

Step 6: Test User Access

Run as: Each demo user role

```
sql
```

```

-- Test as analyst
USE ROLE healthcare_analyst_role;
USE DATABASE healthcare_db;
USE SCHEMA healthcare;
SELECT COUNT(*) FROM claim;

-- Test as power analyst
USE ROLE healthcare_power_analyst_role;
CREATE TEMP TABLE test_temp AS SELECT * FROM claim LIMIT 10;
DROP TABLE test_temp;

-- Test as admin
USE ROLE healthcare_demo_admin_role;
SELECT * FROM demo_audit_log;

```

Post-Installation Tasks

Configure Monitoring

```

sql

-- Enable query history tracking
ALTER SESSION SET USE_CACHED_RESULT = FALSE;

-- Monitor credit usage
SELECT * FROM SNOWFLAKE.ACCOUNT_USAGE.WAREHOUSE_METERING_HISTORY
WHERE WAREHOUSE_NAME = 'HEALTHCARE_DEMO_WH'
ORDER BY START_TIME DESC;

```

Set Up Scheduled Maintenance

```

sql

-- Create task for regular statistics update (optional)
CREATE TASK IF NOT EXISTS update_table_stats
  WAREHOUSE = healthcare_demo_wh
  SCHEDULE = 'USING CRON 0 2 * * SUN America/Denver'
AS
  CALL SYSTEM$GATHER_STATS('healthcare_db.healthcare');

```

Configure Backup (Optional)

```

sql

```

```
-- Enable Time Travel for recovery
```

```
ALTER DATABASE healthcare_db SET DATA_RETENTION_TIME_IN_DAYS = 7;
```

```
-- Create clone for backup
```

```
CREATE DATABASE healthcare_db_backup CLONE healthcare_db;
```

Troubleshooting

Common Issues and Solutions

Issue: Foreign key constraint violations during data load

```
sql
```

```
-- Temporarily disable constraints
```

```
ALTER TABLE claim DROP CONSTRAINT <constraint_name>;
```

```
-- Load data
```

```
-- Re-enable constraints
```

```
ALTER TABLE claim ADD CONSTRAINT <constraint_name>
```

```
FOREIGN KEY (patient_id) REFERENCES patient(patient_id);
```

Issue: Insufficient privileges error

```
sql
```

```
-- Check current role
```

```
SELECT CURRENT_ROLE();
```

```
-- Switch to appropriate role
```

```
USE ROLE SYSADMIN;
```

Issue: Warehouse suspended due to resource monitor

```
sql
```

```
-- Check monitor status
```

```
SHOW RESOURCE MONITORS;
```

```
-- Temporarily increase limit or resume warehouse
```

```
ALTER RESOURCE MONITOR healthcare_demo_monitor
```

```
SET CREDIT_QUOTA = 150;
```

```
ALTER WAREHOUSE healthcare_demo_wh RESUME;
```

Issue: View creation fails due to missing columns

```
sql
```

```
-- Verify base table structure
```

```
DESC TABLE claim;
```

```
-- Check for column name case sensitivity
```

```
SELECT * FROM claim LIMIT 1;
```

Validation Checklist

- ☐ All 17 tables created and populated
- ☐ All 15 views created without errors
- ☐ 6 demo users can login successfully
- ☐ Analyst role can SELECT from tables/views
- ☐ Power analyst can create temp tables
- ☐ Resource monitor active and configured
- ☐ Warehouse auto-suspend working (5 min idle)
- ☐ Sample queries return expected results

Rollback Procedure

If installation fails:

```
sql
```

```
-- Drop in reverse order
```

```
DROP DATABASE IF EXISTS healthcare_db CASCADE;
```

```
DROP WAREHOUSE IF EXISTS healthcare_demo_wh;
```

```
DROP ROLE IF EXISTS healthcare_demo_admin_role;
```

```
DROP ROLE IF EXISTS healthcare_report_writer_role;
```

```
DROP ROLE IF EXISTS healthcare_power_analyst_role;
```

```
DROP ROLE IF EXISTS healthcare_data_steward_role;
```

```
DROP ROLE IF EXISTS healthcare_analyst_role;
```

```
DROP USER IF EXISTS demo_analyst_1;
```

```
-- Continue for all demo users...
```

```
DROP RESOURCE MONITOR IF EXISTS healthcare_demo_monitor;
```

Performance Optimization

Recommended Settings

```
sql
```



```
-- Set warehouse size based on usage
ALTER WAREHOUSE healthcare_demo_wh
SET WAREHOUSE_SIZE = 'MEDIUM' -- For larger demos

-- Enable query acceleration
ALTER WAREHOUSE healthcare_demo_wh
SET ENABLE_QUERY_ACCELERATION = TRUE;

-- Set appropriate clustering keys
ALTER TABLE claim CLUSTER BY (submission_date, claim_status);
ALTER TABLE claim_line_item CLUSTER BY (claim_id);
```

Security Hardening

For production use:

```
sql

-- Enable MFA for admin roles
ALTER USER demo_admin SET MINS_TO_BYPASS_MFA = 0;

-- Set session timeout
ALTER USER demo_analyst_1 SET SESSION_IDLE_TIMEOUT_MINS = 30;

-- Enable audit logging
CREATE OR REPLACE TABLE access_log AS
SELECT * FROM SNOWFLAKE.ACCOUNT_USAGE.QUERY_HISTORY
WHERE DATABASE_NAME = 'HEALTHCARE_DB';
```

Contact Information

- **Installation Issues:** Contact DBA team
- **Access Requests:** Submit via ticketing system
- **Performance Issues:** Monitor warehouse utilization first
- **Data Issues:** Check demo_audit_log table

Data Model Architecture

Logical Sub-Models

The healthcare claims database is organized into six logical sub-models, each representing a distinct business domain:

1. Patient Management Sub-Model

Tables: patient, patient_insurance, insurance_plan

Purpose: Core patient demographics and coverage information

Key Relationships: Patient → Patient_Insurance → Insurance_Plan

Business Owner: Enrollment/Eligibility Team

2. Provider Network Sub-Model

Tables: provider, facility, fee_schedule

Purpose: Healthcare delivery network and contracted rates

Key Relationships: Provider/Facility → Fee_Schedule → Insurance_Plan

Business Owner: Network Management Team

3. Clinical/Medical Sub-Model

Tables: encounter, diagnosis, procedure, prior_authorization

Purpose: Medical events and services that generate claims

Key Relationships: Encounter → Diagnosis/Procedure, Prior_Auth → Procedure

Business Owner: Clinical Operations Team

4. Claims Processing Sub-Model

Tables: claim, claim_line_item

Purpose: Core billing and adjudication transactions

Key Relationships: Claim → Claim_Line_Item

Business Owner: Claims Processing Team

5. Financial/Payment Sub-Model

Tables: payment, payment_adjustment, coordination_of_benefits

Purpose: Post-adjudication financial transactions

Key Relationships: Payment → Payment_Adjustment, COB → Primary/Secondary Claims

Business Owner: Revenue Cycle Team

6. Denial & Appeals Sub-Model

Tables: denial, appeal

Purpose: Exception handling and recovery workflow

Key Relationships: Denial → Appeal (multi-level)

Business Owner: Denial Management Team

Sub-Model Interfaces

Key connection points between sub-models:

- Clinical → Claims: via encounter_id
- Claims → Financial: via claim_id
- Patient Management → Claims: via patient_id and plan_id
- Provider Network → Claims: via provider_id and facility_id
- Claims → Denial & Appeals: via claim_id

Implementation Considerations

For modular deployment:

- Each sub-model can be implemented as a separate schema
- Use database links or views for cross-model queries
- Consider separate warehouses for different teams
- Implement row-level security based on sub-model ownership

For microservices architecture:

- Each sub-model maps to a bounded context
- Define clear API contracts at interface points
- Consider event-driven updates between models
- Implement separate data marts for each domain

Script Execution Order Summary

1. `snowflake_security_setup.sql` (SECURITYADMIN)
2. `healthcare_claims_ddl.sql` (SYSADMIN) - Final schema
3. `raw_schema_ddl.sql` (SYSADMIN) - Optional: RAW landing zone for ETL
4. Data Loading - Choose one approach:
 - **Option A:** `healthcare_sample_data.sql` (direct to final tables)
 - **Option B:** CSV → RAW → dbt → final tables
5. `healthcare_reporting_views.sql` (healthcare_demo_admin_role)

Total estimated installation time: ~10-15 minutes

ETL Architecture Decision

Direct Load (Option A):

- Use for quick demos
- Run `healthcare_sample_data.sql`
- Data goes straight to final tables

ETL Pipeline (Option B):

- Use to demonstrate modern data engineering
- CSV → RAW schema (no constraints) → dbt transforms → HEALTHCARE schema
- Benefits: Data quality checks, lineage tracking, repeatability

RAW Schema Considerations

Why RAW Schema:

- Captures source data exactly as received
- No type conversion failures during initial load
- Audit trail via `_loaded_at` and `_file_name`
- Allows incremental loading and reprocessing

CSV Loading to RAW:

```
sql

-- No foreign key order required for RAW
COPY INTO raw.patients FROM @raw.csv_stage/patients.csv;
COPY INTO raw.claims FROM @raw.csv_stage/claims.csv;
-- Order doesn't matter - no constraints
```

dbt Transformation Example:

```
sql
```

```
-- models/staging/stg_claims.sql
```

```
SELECT
```

```
  claim_id::VARCHAR(50) as claim_id,
```

```
  TRY_TO_DATE(service_date_from) as service_date_from,
```

```
  TRY_TO_DECIMAL(total_charge_amount, 12, 2) as total_charge_amount,
```

```
  -- Data quality checks
```

```
  IFF(claim_status IN ('Approved','Denied','Pending'),
```

```
    claim_status, 'Unknown') as claim_status
```

```
FROM {{ source('raw', 'claims') }}
```

```
WHERE claim_id IS NOT NULL
```